

NLP course 2023

Homework 1

Event Detection

Prof. Roberto Navigli

Teaching assistants:

Edoardo Barba, Tommaso Bonomo,
Karim Ghonim, Giuliano Martinelli,
Francesco Molfese, Stefano Perrella,
Lorenzo Proietti



SAPIENZA
UNIVERSITÀ DI ROMA

SAPIENZA
NLP



Event Detection

An introduction

What is an Event?

- An **event** is a specific occurrence involving participants.
- Each event is labeled with an **event type**.
- An **event trigger** is the **key word or phrase** in an event mention that most clearly expresses the event occurrence.

Dante wrote the divine comedy

EVENT

Event Detection

- Event Detection (ED) is the task of **locating** event triggers and **classifying** event types.
 - It is a sequence labeling task, i.e. every token must be tagged with some label.
 - 5 labels:

SEN: Sentiment

SCEN: Scenario

CHA: Change

POSS: Possession

ACT: Action

Event Detection

Class examples

Protests **ended** after the government **promised** to permanently suspend the project.

Five people **died** in an aviation accident.

Dante **wrote** the divine comedy.

The tornado **attained** category 2 strength.

Dataset

The Dataset

- Each sample is a sentence annotated with **event triggers** and their **classes**.
- Data splits:
 - **20,000** training examples (`train.jsonl`)
 - **2,000** development examples (`dev.jsonl`)
 - **2,000** test examples (`test.jsonl`)
- An event trigger spans across one or more words:
 - This dataset adopts the BIO format to represent event trigger spans.

What is the BIO format?

- It is usually used to label spans of text through token-level tagging.
- **BIO:**
 - **B-** indicates that the tag is the **beginning** of a span
 - **I-** indicates that the tag is **inside** a span
 - **O-** indicates that a token is **outside** of any span

| | | |
|-------------|---------|-------------------|
| | The | 0 |
| ACT | meeting | B-Action |
| | will | 0 |
| SCEN | take | B-Scenario |
| | place | I-Scenario |
| | | 0 |

Dataset format

- JSONL, i.e. each line is a JSON object with three fields:
 - **idx**: progressive index of the sample, unique to each split.
 - **tokens**: the pre-tokenized sentence, i.e. a list of tokens that make up the sentence.
 - **labels**: a list with the same length of **tokens**, with the corresponding label for each token

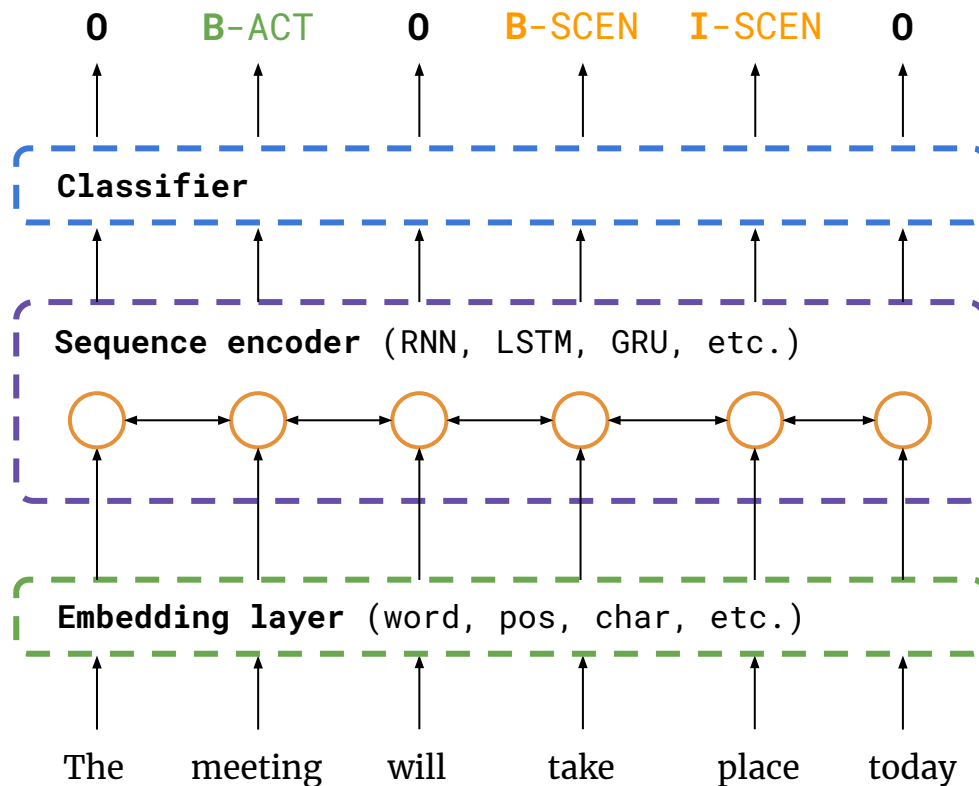
```
{  
  "idx": 10,  
  "tokens": ["The", "meeting", "will", "take", "place", "today"],  
  "labels": ["0", "B-ACTION", "0", "B-SCENARIO", "I-SCENARIO", "0"]  
}
```

Approaches

Model: possible approaches

Note:

Transformer-based
architectures are not allowed



Submission

What you will receive

- We will provide you with a folder organized as follows (some files are omitted):

- nlp2023-hw1/
 - data/
 - hw1/
 - model.py
 - **stud/**
 - **model/**
 - **requirements.txt**
 - test.sh

- You are allowed to edit only the items in bold!

What you will receive

- We will evaluate your work using Docker
 - You should be fine even if you don't know anything about it
- If **test.sh** runs on your side, it will run on ours as well
 - Just keep in mind: do not change any file but those we marked in bold as editable in the previous slide
- Additionally, we wrote a **README.md** to get you everything up and running
- You can find the code repository [here](#)!

What we expect from you

- The zip folder we gave you (but populated :))
- Put your training code (if you used Colab, download the notebook .ipynb and place it) in **hw1/stud/**
- If you use any additional library, modify the **requirements.txt** file as needed (click [here](#) for info)
- Use the data (train, dev and test) in the data folder
 - use each file as defined in the **standard ML conventions** (*train for training, dev for model selection and test for final testing of the model*)

What we expect from you

- Put everything your model needs (vocabulary, weights, ...) inside the **model1/** folder, and be sure to properly load them in your model
- In **hw1/stud/implementation.py** implement the **StudentModel** class
 - Load your model and use it in the **predict** method
 - You must respect the signature of the predict method!
 - You can add other methods (i.e. the constructor)
- In **hw1/stud/implementation.py** implement the **build_model** function
 - It should initialize your **StudentModel** class.

What we expect from you

- Use **test.sh** to check that everything works
- Add your **report.pdf** to the folder (yes, export it in PDF even if you are using Word!)
- Name the zip folder **lastname_studentid_hw1.zip**:
 - Ex: Luigi D'Andrea will submit a file named **dandrea_1234567_hw1.zip**
 - If you are unsure which name to put, use the one in your institutional email account

Submission Instructions

Link sharing

☐ **On - Public on the web**
Anyone on the Internet can find and access. No sign-in required.

☒ **On - Anyone with the link**
Anyone who has the link can access. No sign-in required.

☐ **On - Dipartimento di Informatica - Univ. La Sapienza**
Anyone at Dipartimento di Informatica - Univ. La Sapienza can find and access.

☐ **On - Anyone at Dipartimento di Informatica - Univ. La Sapienza with the link**
Anyone at Dipartimento di Informatica - Univ. La Sapienza who has the link can access.

☐ **Off - Specific people**
Shared with specific people.

Access: Anyone (no sign-in required) [Can view](#) ▼

Viewers of this file can see comments and suggestions. [Learn more](#)

Note: Items with any link sharing option can still be published to the web. [Learn more](#)

[Save](#) [Cancel](#) [Learn more about link sharing](#)

- Upload the zip on your **institutional** Drive and make it **link-shareable** and **public** to anyone (an automatic script will download it).
- Make sure it is accessible via an incognito page of your browser!
- Do **NOT modify** the folder structure
- You have to submit the homework through the [submission form](#) on Google Classroom. You will be asked to fill a form with the requested information and the **link** to the zip you uploaded on Drive.

Evaluation

Evaluation

- Use the **validation split** to select the **best model/hyperparameters** configuration
- Use the **test split** to evaluate your model and **estimate its performance**
- The final evaluation will be conducted on a **SECRET** test set
- The evaluation metric will be the **macro F1-score** obtained comparing your model's predictions with our golden labels
- We will use [segeval](#), a common Python package used to obtain **token-level classification metrics**, one of which is the F1-score

Evaluation

F1 explanation

- The **F1-score** for a single class C is defined as:

$$F_1 = 2 \cdot \frac{\text{recall} \times \text{precision}}{\text{recall} + \text{precision}}$$
$$\text{recall} = \frac{\text{"number of tokens labelled as } C \text{ that are truly part of class } C\text{"}}{\text{"number of all tokens that are truly part of } C\text{"}}$$
$$\text{precision} = \frac{\text{"number of tokens labelled as } C \text{ that are truly part of class } C\text{"}}{\text{"number of all tokens labelled as } C\text{"}}$$

- The **macro F1-score** is simply the average of the F1-score for each class
- We use “strict” mode in `seqeval`, i.e. tags must adhere correctly to the BIO format




B-SCEN I-SCEN 0



I-SCEN I-SCEN 0

Evaluation


We will take into account the following criteria:



Report: dos and don'ts

- **ACL 2023 paper template**
 - Freely available: [LaTeX](#), [Word](#) or [Overleaf](#)
 - You can use either the LaTeX or the Word template, your choice
 - **DO NOT MODIFY** the template (margins, spacing, font size)
 - Use the non-anonymous flag, so you can enter your name
- **Max 2 pages**
 - For the report, including title, subtitles, etc.
 - This is a **STRICT RULE!**
- **Unlimited extra pages for images, tables and references**
 - Be sure to **include** and properly **comment** a [confusion matrix](#), visualized as heat map
 - Every image and table must have a caption (don't abuse them please :))
 - Tables and images must be referenced in the report

Confusion Matrix



Report: what you are expected to do



We expect a good report to be:

- **Readable and understandable**
 - We will not give penalties for English errors, but we expect the report to follow a clear flow. We don't want to read just a sequence of statements on what you did without showing the reasoning behind your choices
- **Well-structured and organized**
 - Take inspiration from the many papers available online and organize your report in well-defined sections (e.g. method, setup, experiments, results...)

Report: what you are not expected to do



We expect a good report **NOT** to include:

- Unnecessary **task** or **dataset descriptions**
 - just focus on your solution to the problem
- **Code** copy-paste
 - Your code should be self-explanatory, so no need to show it in the report. You can add **pseudocode** to show some particular algorithm, but **no code or screenshots**, please!

Report: what you are not expected to do



We expect a good report **NOT** to include:

- **Unnecessary low-level implementation details**
 - Avoid any **low-level implementation/technical details** like “I used a dictionary to store these values”, “I had to use configuration X to solve this exception”, “I could not use Y because there was a dependency issue with Z”, etc.
 - Instead, **we are interested in high-level abstractions/strategies** you decide to use to tackle the homework, as well as the **intuitions behind your choices**.
E.g. use and description of a particular model, explanation of how and why an architecture works, etc.

Application: what you are expected to do



Your project should conform to the following rules:

- You **MUST** use PyTorch.
 - TensorFlow and other deep learning frameworks are **NOT** allowed.
 - PyTorch Lightning is **NOT** allowed (at this stage)
- **Frameworks** that use PyTorch (e.g. AllenNLP, torchtext...) are **NOT** allowed.
- Libraries (such as tqdm, sklearn, NLTK) are fine, but since the line between a framework and a library is sometimes blurred, please ask in the Google Classroom group before using any external library: **any other library MUST be agreed with the TAs.**

Application: what you are not expected to do



Your project should conform to the following rules:

- **You are not allowed** to use tools/architectures that have not been explained yet in the course, in particular:
 - word embeddings (Word2Vec, GloVe, etc.) **are allowed**,
 - contextualized word embeddings (ELMo, etc.) **are NOT allowed**,
 - Transformer-based models (BERT, BART, RoBERTa, etc.) **are NOT allowed**.
- For any doubt, please ask the TAs on Google Classroom.
- **Comment** your code, please!

Quantitative Results

We will evaluate the **performance of your model** on a SECRET test set.

You can get **from 0 to 6** points according to the following **thresholds**:

- $F1 < 0.15$ \Rightarrow FAIL
- $0.15 < F1 < 0.30$ \Rightarrow 0
- $0.30 < F1 < T2$ \Rightarrow 1
- $T2 < F1 < T3$ \Rightarrow 2
- $T3 < F1 < T4$ \Rightarrow 3
- $T4 < F1 < T5$ \Rightarrow 4
- $T5 < F1 < T6$ \Rightarrow 5
- $F1 > T6$ \Rightarrow 6

Quantitative Results

We will evaluate the **performance of your model** on a SECRET test set.

You can get **from 0 to 6** points according to the following **thresholds**:

- $F1 < 0.15$ \Rightarrow FAIL
- $0.15 < F1 < 0.30$ \Rightarrow 0
- $0.30 < F1 < T2$ \Rightarrow 1
- $T2 < F1 < T3$ \Rightarrow 2
- $T3 < F1 < T4$ \Rightarrow 3
- $T4 < F1 < T5$ \Rightarrow 4
- $T5 < F1 < T6$ \Rightarrow 5
- $F1 > T6$ \Rightarrow 6

Thresholds will be defined
based on an internal reference
model and the **normalized
distribution of YOUR scores!**

Extras

You can achieve **up to 5 points with some extras!**

An “extra” is whatever you decide to add to your model to make it better. For instance:

- use of pre-trained embeddings, PoS embeddings, N-grams, char embeddings, CRF, ...
- use of NLP best practices,
- comparative analysis of results in your report,
- informative plots in your report,
- **new ideas** (including using other approaches in a clever way)

and more, according to internal baselines. Don't forget to **explain your choices** in the report! Extras that are not explained in the report will not be considered for evaluation.

Evaluation

- `test.sh` is identical to what we will be using
- **If it does not run on your side, we will not correct your homework**
- Note that, if you use **any kind of hard-coded paths**, this script won't work
- Use [paths relative](#) to the project root folder, e.g.:
 - **NO:** `/home/pincopallino/my_folder/model/weights.pt`
 - **OK:** `model/weights.pt`

Warnings

Things you should be aware of



SAPIENZA
NLP



Please be aware that

This is an **individual exercise**! Collaboration among the students is **not** allowed.

We will check for **plagiarism** both manually and automatically.

It is **not allowed** to:

- Copy from other students.
- Share your code with other students.
- Use ChatGPT or similar systems for report writing.
- Copy from online resources (StackOverflow, GitHub, Medium, Kaggle and so on).

You are also allowed to use the **SOME** parts of the presented class notebooks. However, you **MUST** explicitly specify these parts in your code comments.

WARNING if a notebook uses torchtext or pytorch-lightning you **cannot** use it

Data policy

- For your experiments, use **ONLY** the provided data (train, dev and test) in the data folder; use each file as defined in the standard ML conventions (train for training, dev for model selection and test for testing).
- If you train it on dev or test set, it will be a **FAIL**.

Tips



SAPIENZA
NLP



A few tips to organize your work:

- **Start as soon as possible!**
 - Training a neural network requires time, possibly hours, depending on your hardware
- **Start small!**
 - If you don't get decent results with a very simple neural network, there is a good chance that adding other things won't make your model perform better
- **Leave some time for hyperparameter tuning!**
 - Sometimes good hyperparameter combinations can do wonders for your neural network
- **Use Google [Colab](#) (free GPUs!)**

Deadline

When to deliver what



Deadline

Submission date: **May 1st, 2023 (Monday)**
23:59:59 Italian time (UTC + 1)

Submit the homework through the submission form on Google Classroom. You have to fill the form with the requested information and a link to the zip folder of the homework on Google Drive.

Late submission policy

Late Submission date: until **May 7th, 2023 (Wednesday)**
23:59:59 Italian time (UTC + 1)

1 point penalty will be applied for each day of delay, e.g.:

- A student delivers their homework on May 4th -> max possible grade $35 - 3 = 32$
- A student delivers their homework on May 8th -> **FAIL!**

Awards

Get a Sapienza NLP™ t-shirt




SAPIENZA
NLP



Win a Sapienza NLP t-shirt!

We will hand out amazing Sapienza NLP t-shirts to the **overall top-5** students!

The final ranking will be computed according to the scores on our **secret** test set.



That's not all

If your work is novel, interesting and original, we will gladly invite you to work together with us to extend on a fully-fledged paper for **TOP-TIER INTERNATIONAL CONFERENCE!**

Just over the last 12 months, the Sapienza NLP group published more than a dozen of papers!

Questions?

If you have a question that may interest your colleagues, **please ask it on Google Classroom.**

Otherwise, for personal or other questions, email **ALL** of us (but please, only reach for things that can't be asked on the Google Classroom).

Our emails are:

{bonomo, ghonim, martinelli, molfese, perrella, lproietti}@diag.uniroma1.it

Good Luck!!

